

Topic : Parallelizing QuickSort using CilkPlus.

Students : Hao Hu (3576231), Abhijit Taware(3590004).

Supervisor : Eric Aubanel (Parallel Programming course CS 6025).

QuickSort is efficient sorting algorithm and is mainly a divide and conquer algorithm. The algorithm divides the list of elements in two parts and recurse on each half. The partitioning is done by selecting a pivot element. All the elements smaller than or equal to the pivot goes to one partition and rest goes to the other. This division eventually ensures that the array is sorted. The pseudo code can be as follows -

```
algorithm quicksort(A, lo, hi) is
  if lo < hi then
    p := partition(A, lo, hi)
    quicksort(A, lo, p - 1)
    quicksort(A, p + 1, hi)
```

As a first deliverable we are focusing on implementing the partition efficiently. Parallelizing the recursive calls should be easy using `cilk_spawn` and `cilk_sync`.

E.g.

A = [7, 4, 5, 3, 19, 10, 11, 9, 1]

If we select first element i.e. 7 as pivot then the the two partitions will be as follows -
[7, 4, 5, 3, 1] and [19, 10, 11, 9]

Partitioning it using sequential algorithm has $O(n)$ complexity and would take longer for a very huge value of n . We can use prefix sum to partition the array in parallel. The trick is to keep additional indicator array **B** which encodes information about whether particular element in array **A** is smaller or not. Array **B** will have value 1 for if $A[i] \leq \text{pivot}$, 0 otherwise. Now doing prefix sum of **B** gives us way to partition array **A**. The array **B** would be as follows for above example -

B = [1,1,1,1,0,0,0,0,1] . It prefix sum would be \rightarrow **B_sum** = [0,1,2,3,3,3,3,3,4]

The **B_sum** gives us indices for first half of the partition, i.e. the left half would be -

L = [7, 4, 5, 3, 1] , note we have to use indices from **B_sum** for the elements which have value 1 in array **B**.

As a part of first deliverable we have to implement exclusive prefix sum. We have done it using Blelloch scan algorithm mentioned in the book and at the following link -
<https://www.cs.cmu.edu/~guyb/papers/Ble93.pdf>

This scan algorithm works in two phases viz. upsweep and downsweep. The upsweep is a regular reduction as given below.

```

Procedure up-sweep(a)
for d from 0 to (lg n) - 1
  in parallel for i from 0 to n - 1 by 2d+1
     $a[i + 2^{d+1} - 1] \leftarrow a[i + 2^d - 1] + a[i + 2^{d+1} - 1]$ 

```

If we have array $A = [1, 2, 3, 4, 5, 6, 7, 8]$, following steps show the complete up-sweep

1	2	3	4	5	6	7	8
	3		7		11		15
		10				26	
						36	

Dow sweep starts by setting an identity element i.e. last element is set to zero. The pseudo code as follow.

```

procedure down-sweep(a)
   $a[n - 1] \leftarrow 0$ 
  for d from (lg n) - 1 downto 0
    in parallel for i from 0 to n - 1 by 2d+1
       $t \leftarrow a[i + 2^d - 1]$ 
       $a[i + 2^d - 1] \leftarrow a[i + 2^{d+1} - 1]$ 
       $a[i + 2^{d+1} - 1] \leftarrow t + a[i + 2^{d+1} - 1]$ 

```

1	2	3	4	5	6	7	8	
	3		7		11		15	
		10				26		
						36		
			10			0		← Setting identity element and using values from upsweep.
			0			10		← setting left and right children
	3		0		11		10	
	0		3		10		21	
1	0	3	3	5	10	7	21	
0	1	3	6	10	15	21	28	

We have used Cilk plus for this implementation. Instead of using `cilk_for`, we are relying on array notations in cilk plus to vectorize the code. Detailed implementation can be found in the source code provided with the report.

Above implementation assumes the array size can be expressed as power of two. However in practice we encounter different array sizes and hence need to modify the algorithm to handle this case as well. We are proposing two different implementations for handling arrays of different sizes -

1. Zero filling the array to match next smallest value which is power of 2
2. Implement the Blleloch scan from from left to right and pretend there are more elements in the end to fulfill power of 2 requirement.

1. Zero filling the array -

Suppose we have an array with 5 elements as follows -

1 2 3 4 5 , we extend it to an array with 8 elements and zero fill the additional elements.
1 2 3 4 5 0 0 0

Now the blleloch scan can proceeds as usual. The downside is we have allocate space for extra elements.

2. Reverse Blleloch scan -

In an attempt to avoid zero filling, we have tried a different approach to Blleloch scan which can be described as follows.

- Reverse the original array
- Start up sweep by adding element on right to element on left
 - $A[\text{left}] += A[\text{right}]$
- Due to above step we don't need additional space. Whenever the 'right' pointer goes past end of array, we return a zero value for that element.
- Downsweep also happens in reverse direction.
- However we get the prefix sum array in reverse direction and we have to reverse the final result again.

Following example shows this simple trick

14	9	3	11	8	← original array
8	11	3	9	14	← reversed array
19		12	14	0 0 0	← $A[\text{left}] += A[\text{right}]$ and dummy zero
31			14	0	
45					
0	11	12	9	14	← setting identity element & reusing results from upsweep
14				0	
14		12	9	0	
26	11	14	9	0	
37	26	23	14	0	← this is reversed prefix sum
0	14	23	26	37	← reverse again to get the prefix sum

Following sections gives results for both the implementations. For reversed Blleloch scan, we have not included the timing of reverse in the result.

The last reverse operation can avoided by manipulating index.

E.g. $\text{index} = (\text{index} + \text{size} - 1) \% \text{size}$

Results :-

The code is tested on compute node 0, which 24 core 3.5Hz Intel Xeon machine.

\$ uname -a

```
Linux compute-0-0.local 2.6.32-431.11.2.el6.x86_64 #1 SMP Tue Mar 25 19:59:55 UTC
2014 x86_64 x86_64 x86_64 GNU/Linux
```

\$ cat /proc/cpuinfo | grep processor | wc -l

```
24
```

Code can be compiled using standard make command on compute-0-0 node in lab.
Steps to compile and run.

- make ← compiles prefix and r_prefix
- make test ← runs test for different array values for both implementations.

prefix ⇐ zero filled array implementation.

r_prefix ⇐ reversed blelloch scan implementation.

-bash-4.1\$ make

```
gcc -fcilkplus -O2 -ftree-vectorize -fopt-info-vec-optimized -lrt -o prefix prefix.c
gcc -fcilkplus -O2 -ftree-vectorize -fopt-info-vec-optimized -lrt -o r_prefix r_prefix.c
r_prefix.c:20:13: note: loop vectorized
```

-bash-4.1\$ make test

```
bash test.sh
```

```
-----
Testing standard Blelloch prefix sum implementation.
-----
```

Array size : 32

sequential prefix sum time in seconds: 0.000008

parallel prefix sum time in seconds : 0.000004

correct result.

Array size : 1023

sequential prefix sum time in seconds: 0.000015

parallel prefix sum time in seconds : 0.000013

correct result.

Array size : 32767

sequential prefix sum time in seconds: 0.000248

parallel prefix sum time in seconds : 0.000444

correct result.

Array size : 33554431

sequential prefix sum time in seconds: 0.081702

parallel prefix sum time in seconds : 0.429251

correct result.

Testing reversed Blelloch prefix sum implementation.

Array size : 32

sequential prefix sum time in seconds: 0.000009

parallel prefix sum time in seconds : 0.000003

correct result.

Array size : 1023

sequential prefix sum time in seconds: 0.000014

parallel prefix sum time in seconds : 0.000013

correct result.

Array size : 32767

sequential prefix sum time in seconds: 0.000231

parallel prefix sum time in seconds : 0.000555

correct result.

Array size : 33554431

sequential prefix sum time in seconds: 0.082191

parallel prefix sum time in seconds : 0.436850

correct result.

We have been able to get the algorithm working correctly. The numbers are really bad with parallel scan and we are investigating the reasons for the same.

References :

Book : Elements of Parallel Computing by Prof. Eric Aubanel

Publisher : Chapman and Hall/CRC ISBN 9781498727891

Paper : Prefix Sums and Their Applications

Guy E. Blelloch Carnegie Mellon University